


```
library(caret)
```

```
rladies_global %>%  
  filter(city == 'Leuven')
```



Classification



This book
is one of the best machine
learning books out there.

It's also free.

<http://www-bcf.usc.edu/~gareth/ISL/ISLR%20First%20Printing.pdf>



We have a labelled dataset.

Outcome A

- Customer
- Disease status
- Water condition

**Outcome
B**



We have a labelled dataset.

Outcome A



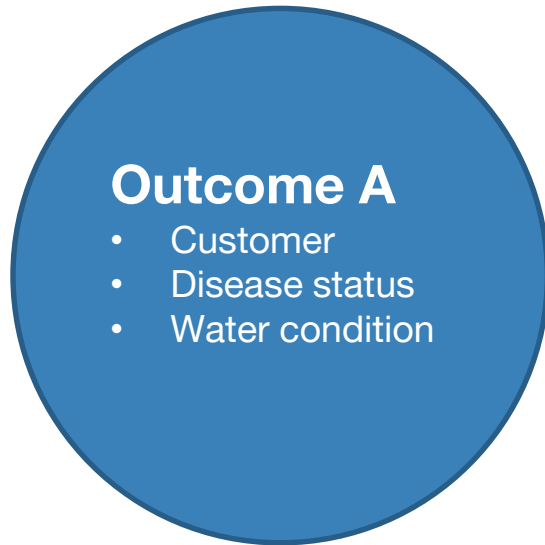
- Customer
- Disease status
- Water condition

Outcome B

Can we predict whether a new data point is a **circle** or a **triangle**?



We have a labelled dataset.




Let's see how these algorithms work.




K Nearest Neighbours

Labels are assigned
to known datapoints.



Labels are assigned to known datapoints.


New point is placed within **known space**.



Labels are assigned to known datapoints.

New point is placed within **known space**.

Find **K neighbours** (here, $K = 3$).




Labels are assigned to known datapoints.

New point is placed within **known space**.

Find **K neighbours** (here, $K = 3$).

All neighbours are circles -> **new point** is a circle.





Logistic Regression




The **logistic function** returns probabilities between 1 and 0.

$$\sigma(t) = \frac{1}{1 + e^{-t}}$$

The **logistic function** returns probabilities between 1 and 0.

$$\sigma(t) = \frac{1}{1 + e^{-t}}$$






The **logistic function** returns probabilities between 1 and 0.

$$\sigma(t) = \frac{1}{1 + e^{-t}}$$

We cannot perform linear regression on this function, because it is non-linear.




The **logistic function** returns probabilities between 1 and 0.

$$\sigma(t) = \frac{1}{1 + e^{-t}}$$

We cannot perform linear regression on this function, because it is non-linear.

With a known t , we can define the **logit function** by taking the inverse of the **logistic function**.






The **logistic function** returns probabilities between 1 and 0.

$$\sigma(t) = \frac{1}{1 + e^{-t}}$$

We cannot perform linear regression on this function, because it is non-linear.

With a known t , we can define the **logit function** by taking the inverse of the **logistic function**.




The **logit function** is linear. Therefore, we can use linear regression to fit a model.



Random Forest

Random forests are like simple **decision trees**.


A **decision tree** finds the best way to split a dataset.



Random forests are like simple **decision trees**.

A **decision tree** finds the best way to split a dataset.

It looks through **all predictors** to find the one that has the smallest prediction error.




Random forests are like simple **decision trees**.

A **decision tree** finds the best way to split a dataset.

It looks through **all predictors** to find the one that has the smallest prediction error.

Then, a new branch is generated.




Random forests are like simple **decision trees**.

A **decision tree** finds the best way to split a dataset.

It looks through **all predictors** to find the one that has the smallest prediction error.

Then, a new branch is generated.

This results in high sensitivity to noise and to overfitting.




Random forests are like simple **decision trees**.

A **decision tree** finds the best way to split a dataset.

It looks through **all predictors** to find the one that has the smallest prediction error.

Then, a new branch is generated.




This results in high sensitivity to noise and to overfitting.






**Why use one tree when
you can have a forest?**

Random forests are made from multiple decision “stumps”.




Random forests are made from multiple decision “stumps”.



Each **stump** is given a training set **randomly sampled with replacement** from the original set.


Random forests are made from multiple decision “stumps”.



Each **stump** is given a training set **randomly sampled with replacement** from the original set. They use one predictor.



Random forests are made from multiple decision “stumps”.



Each **stump** is given a training set **randomly sampled with replacement** from the original set. They use one predictor.

The **random forest** averages the predictions (regression) or takes the majority vote (classification).

Classification with caret



Data preparation



Model training



Model evaluation



Find a dataset

`https://archive.ics.uci.edu/ml/datasets/Cervical+cancer+%28Risk+Factors%29`

Cervical cancer (Risk Factors) Data Set

Fernandes, K., Cardoso, J. S., & Fernandes, J. (2017, June). Transfer learning with partial observability applied to cervical cancer screening. In *Iberian conference on pattern recognition and image analysis* (pp. 243-250). Springer, Cham.

```
library(datasets)
data(iris)
```



Model training



Replace ? with NA

`read.csv`



**Split up your data in a
training and a test set**


`createDataPartition`



Train your model

`train`

Prevent overfitting!



Model testing



Test model

`predict`

Evaluate predictions

`confusionMatrix`
`LOOCV`
`Boot`
`cv`

Improve model

`kappa`
`trControl`



Unlabeled data

**Upclass
RSSL**